



Analyzing Structured Data in Python

WEEK 1: PANDAS

14 March 2022

COURSE STRUCTURE

Anticipate about ~7 hours/week

- 2 course meetings per week
 - 10:00 - 11:00 AM BST Mondays
 - 10:00 - 11:00 AM BST Fridays
- 1 assignment per week ~2 hours
- Office hours on Wednesdays for 30 minutes per participant
- Independent learning ~2 hours

Teams for introductions, meetings, office hours, questions, files

COURSE TOPICS

Week 1: Pandas for CSV data

Week 2: ElementTree for XML data

Please let me know if there are specific topics you'd like to cover!

MORE PYTHON COURSES WITH THE CDCS

Machine Learning with Python

23 March - 13 April (2 days a week)

Text Analysis with Python's NLTK Library

11 April - 22 April (2 days a week)

The CDCS has many resources on its website for courses and self-guided learning!

INTRODUCTIONS

Name

Have you programmed before?
If yes, in what language?

Why would you like to learn Python?

FOR PARTICIPANTS

- Introduce material for you to review in greater depth on your own
- I'll direct you to further resources if you'd like to go beyond material covered in each week's assignment
- Course meetings won't be recorded
 - Three strike policy
 - Please let me know in advance if you cannot attend!
- Office hours: questions about assignments, your own projects
 - *Chat with me on Teams to schedule!*

FOR PARTICIPANTS

Python version 3

Jupyter Notebooks - options:

- A. Use Noteable: <https://www.ed.ac.uk/information-services/learning-technology/noteable/accessing-noteable>
- B. Use GoogleColab: <https://colab.research.google.com>
- C. Install to your computer: <https://jupyter.org/install>

LET'S CODE!

FURTHER RESOURCES

- Noteable User Guide: https://noteable.edina.ac.uk/user_guide/#hide_ge_7
- Jupyter Notebooks, Noteable: <https://github.com/edina/Exemplars2020/blob/master/TeachingDocs/Tutorials/UsingNoteableBeginner.ipynb>
- Jupyter Notebooks: <https://glam-workbench.github.io/getting-started/>
- Python: <https://programminghistorian.org/en/lessons/introduction-and-installation>

PANDAS

Pandas is built on NumPy

Pandas-specific data structures:

1. Series
2. DataFrame

Reference: https://pandas.pydata.org/pandas-docs/stable/user_guide/10min.html

PANDAS

DataFrames are like tables, storing data in rows and columns

- Text
- Numbers
- Lists, arrays, dictionaries

Each column can have a unique data type, or **dtype**

PANDAS

DataFrames can be created from data in numerous formats, including lists and arrays...

```
In [5]: dates = pd.date_range('20130101', periods=6)
```

```
In [6]: dates
```

```
Out[6]:
```

```
DatetimeIndex(['2013-01-01', '2013-01-02', '2013-01-03', '2013-01-04',  
              '2013-01-05', '2013-01-06'],  
              dtype='datetime64[ns]', freq='D')
```

```
In [7]: df = pd.DataFrame(np.random.randn(6, 4), index=dates, columns=list('ABCD'))
```

```
In [8]: df
```

```
Out[8]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

PANDAS

...dictionaries...

```
In [9]: df2 = pd.DataFrame({'A': 1.,  
    ....:                  'B': pd.Timestamp('20130102'),  
    ....:                  'C': pd.Series(1, index=list(range(4)), dtype='float32'),  
    ....:                  'D': np.array([3] * 4, dtype='int32'),  
    ....:                  'E': pd.Categorical(["test", "train", "test", "train"]),  
    ....:                  'F': 'foo'})
```

```
In [10]: df2
```

```
Out[10]:
```

	A	B	C	D	E	F
0	1.0	2013-01-02	1.0	3	test	foo
1	1.0	2013-01-02	1.0	3	train	foo
2	1.0	2013-01-02	1.0	3	test	foo
3	1.0	2013-01-02	1.0	3	train	foo

PANDAS

...and CSV files

```
In [142]: pd.read_csv('foo.csv')
```

```
Out[142]:
```

	Unnamed: 0	A	B	C	D
0	2000-01-01	0.350262	0.843315	1.798556	0.782234
1	2000-01-02	-0.586873	0.034907	1.923792	-0.562651
2	2000-01-03	-1.245477	-0.963406	2.269575	-1.612566
3	2000-01-04	-0.252830	-0.498066	3.176886	-1.275581
4	2000-01-05	-1.044057	0.118042	2.768571	0.386039
...
995	2002-09-22	-48.017654	31.474551	69.146374	-47.541670
996	2002-09-23	-47.207912	32.627390	68.505254	-48.828331
997	2002-09-24	-48.907133	31.990402	67.310924	-49.391051
998	2002-09-25	-50.146062	33.716770	67.717434	-49.037577
999	2002-09-26	-49.724318	33.479952	68.108014	-48.822030

```
[1000 rows x 5 columns]
```


PANDAS

If you have a lot of data, you can select a subset of rows to view

```
In [13]: df.head()
```

```
Out[13]:
```

	A	B	C	D
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401

```
In [14]: df.tail(3)
```

```
Out[14]:
```

	A	B	C	D
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-05	-0.424972	0.567020	0.276232	-1.087401
2013-01-06	-0.673690	0.113648	-1.478427	0.524988

LET'S CODE!

PANDAS

There are lots of built in methods to reorganize your data...

```
In [20]: df.T
```

```
Out[20]:
```

	2013-01-01	2013-01-02	2013-01-03	2013-01-04	2013-01-05	2013-01-06
A	0.469112	1.212112	-0.861849	0.721555	-0.424972	-0.673690
B	-0.282863	-0.173215	-2.104569	-0.706771	0.567020	0.113648
C	-1.509059	0.119209	-0.494929	-1.039575	0.276232	-1.478427
D	-1.135632	-1.044236	1.071804	0.271860	-1.087401	0.524988

```
In [21]: df.sort_index(axis=1, ascending=False)
```

```
Out[21]:
```

	D	C	B	A
2013-01-01	-1.135632	-1.509059	-0.282863	0.469112
2013-01-02	-1.044236	0.119209	-0.173215	1.212112
2013-01-03	1.071804	-0.494929	-2.104569	-0.861849
2013-01-04	0.271860	-1.039575	-0.706771	0.721555
2013-01-05	-1.087401	0.276232	0.567020	-0.424972
2013-01-06	0.524988	-1.478427	0.113648	-0.673690

```
In [22]: df.sort_values(by='B')
```

```
Out[22]:
```

	A	B	C	D
2013-01-03	-0.861849	-2.104569	-0.494929	1.071804
2013-01-04	0.721555	-0.706771	-1.039575	0.271860
2013-01-01	0.469112	-0.282863	-1.509059	-1.135632
2013-01-02	1.212112	-0.173215	0.119209	-1.044236
2013-01-06	-0.673690	0.113648	-1.478427	0.524988
2013-01-05	-0.424972	0.567020	0.276232	-1.087401

PANDAS

...calculate statistics...

```
In [61]: df.mean()
```

```
Out[61]:
```

```
A    -0.004474
B    -0.383981
C    -0.687758
D     5.000000
F     3.000000
dtype: float64
```

```
In [62]: df.mean(1)
```

```
Out[62]:
```

```
2013-01-01    0.872735
2013-01-02    1.431621
2013-01-03    0.707731
2013-01-04    1.395042
2013-01-05    1.883656
2013-01-06    1.592306
Freq: D, dtype: float64
```

```
In [19]: df.describe()
```

```
Out[19]:
```

	A	B	C	D
count	6.000000	6.000000	6.000000	6.000000
mean	0.073711	-0.431125	-0.687758	-0.233103
std	0.843157	0.922818	0.779887	0.973118
min	-0.861849	-2.104569	-1.509059	-1.135632
25%	-0.611510	-0.600794	-1.368714	-1.076610
50%	0.022070	-0.228039	-0.767252	-0.386188
75%	0.658444	0.041933	-0.034326	0.461706
max	1.212112	0.567020	0.276232	1.071804

PANDAS

...and group and reshape your data!

```
In [87]: df = pd.DataFrame({'A': ['foo', 'bar', 'foo', 'bar',  
.....:                        'foo', 'bar', 'foo', 'foo'],  
.....:                    'B': ['one', 'one', 'two', 'three',  
.....:                        'two', 'two', 'one', 'three'],  
.....:                    'C': np.random.randn(8),  
.....:                    'D': np.random.randn(8)})
```

```
In [88]: df  
Out[88]:
```

	A	B	C	D
0	foo	one	1.346061	-1.577585
1	bar	one	1.511763	0.396823
2	foo	two	1.627081	-0.105381
3	bar	three	-0.990582	-0.532532
4	foo	two	-0.441652	1.453749
5	bar	two	1.211526	1.208843
6	foo	one	0.268520	-0.080952
7	foo	three	0.024580	-0.264610

```
In [89]: df.groupby('A').sum()  
Out[89]:
```

	C	D
A		
bar	1.732707	1.073134
foo	2.824590	-0.574779

```
In [90]: df.groupby(['A', 'B']).sum()  
Out[90]:
```

		C	D
A	B		
bar	one	1.511763	0.396823
	three	-0.990582	-0.532532
	two	1.211526	1.208843
foo	one	1.614581	-1.658537
	three	0.024580	-0.264610
	two	1.185429	1.348368

LET'S CODE!

ASSIGNMENT

Watch the videos in sections *6. Introduction to Pandas* and *7. Baby Names with Pandas* in the course *Python: Data Analysis (2015)*

<https://www.linkedin.com/learning/python-data-analysis-2015/dataframes-in-pandas?u=50251009>

Follow along in your own Jupyter Notebook!

Find your own CSV file to load and turn into a DataFrame (or create your own). What questions can you ask about it using the methods and functions in Pandas?

An icon consisting of a light blue rounded rectangle. Inside, there is a darker blue rounded rectangle with a white grid pattern, resembling a calendar. Overlapping the bottom right of the calendar is another light blue rounded rectangle containing a dark blue code symbol: a less-than sign, a forward slash, and a greater-than sign (</>).

THANKS EVERYONE!

Next course meeting: Friday, 10:00-11:00 AM BST

Office hours available on Wednesday (30 minutes)

To schedule, please message me on Teams!